

Platform Master Roadmap

M1-M7 | Central truth, classification and AI

SOURCE-GROUNDED PLANNING EDITION

Prepared 06 October 2026. Milestone numbering and scope are preserved. This document is a roadmap, not milestone acceptance or permission to dispatch paid requests.

Current position

- M1: accepted for the recorded snapshot. M2: not accepted. M3-M7: not formally completed under the acceptance sequence.
- The latest supplied preparation history is not evidence of completed fresh validation or production cutover.

Sequence

- M1 - Data Requirements, Ownership & Current-State Map
- M2 - Extraction Reliability
- M3 - Central Document Classification & Attribution
- M4 - Central Processing, Relationships & Domain Records
- M5 - Historical Backfill, Registry Completeness & Validation
- M6 - Database-Driven Tab Migration
- M7 - Stabilization, Dependency Engine & Legacy Cleanup

Data Requirements, Ownership & Current-State Map

Define one owner and one producer for every engineering fact before rebuilding consumers.

Work scope

- Inventory each tab, field, reader, API, database table and manual override.
- Map current processing and identify duplicate interpretation, missing ownership and business-state boundaries.

Deliverables

- Field ownership matrix, current-state map and protected behaviour inventory.

Acceptance gate

- Every important field has a producer, persistence owner and consumer contract.
- Existing engineer decisions and business workflows are explicitly protected.

Boundary and dependency

- M1 is accepted for its recorded reviewed snapshot; later changes need traceable updates.

Extraction Reliability

Produce reliable, source-bound observations before classification or domain promotion.

Work scope

- Read identities, revisions, dates, systems, titles, floors and consultant decisions with correct page/component association.
- Verify BOQ part numbers and quantities; preserve punctuation, row identity and ambiguity.
- Use AI primary reading, targeted verification, confident-output audits and bounded escalation where justified.
- Compare baseline and candidate on reviewed frozen labels; include unread pages, unsupported files, failures and budget stops.

Deliverables

- Versioned evidence, corrected readers, frozen harness, accuracy/coverage/cost reports and fresh validation package.

Acceptance gate

- Fix confirmed defects; meet the governing accuracy and recovery thresholds.
- No unresolved critical false acceptance; demonstrate compatibility and reviewed evidence.
- Measured improvements must survive project/layout concentration checks and independent review.

Boundary and dependency

- M2 remains CHANGES STILL REQUIRED in the supplied history. Merged delivery readiness is not M2 acceptance.
- Do not silently renew budgets, reopen terminal stops or reinterpret held labels after prediction.

Central Document Classification & Attribution

Determine what a document is and whose scope it belongs to, without granting business approval.

Work scope

- Classify document type, system, component, context and project/scope association.
- Keep UNKNOWN and AMBIGUOUS states, source provenance, classifier version and freshness.
- Provide an inspector and evaluate shadow precision/recall and false-supported or false-OUR_SCOPE outcomes.

Deliverables

- Versioned classification records, attribution evidence, inspector and validation report.

Acceptance gate

- Classification is explainable and traceable; unsupported assertions are held.
- Type/scope quality is measured before consumers depend on it.

Boundary and dependency

- Begins after M2 acceptance. Existing Classification V2 code does not establish milestone acceptance.

Central Processing, Relationships & Domain Records

Process centrally and persist engineering truth once for all tabs.

Work scope

- Run background domain processors over stored observations and classification.
- Persist document/page/component relationships, drawing revisions, submittal cycles, comments and replies.
- Separate raw observations from effective engineer-approved values; retain lineage and overrides.
- Build domain APIs and dependency relationships; reading a tab must not reopen source documents.

Deliverables

- Central processing contracts, relationship records, domain APIs, job history and sample reconciled projects.

Acceptance gate

- Background processors are repeatable and recoverable.
- Domain facts have evidence and ownership; updates retain history and do not overwrite approved truth.

Boundary and dependency

- Classification describes the file; domain processing creates structured records under explicit rules.

Historical Backfill, Registry Completeness & Validation

Reconcile legacy project data before switching all consumers to the central model.

Work scope

- Discover complete file registries and remove silent enumeration or record-count limits.
- Backfill and reconcile old facts, associations and relationships in isolated batches.
- Audit before/after changes, missing records, duplicates and preserved manual decisions.

Deliverables

- Complete registries, migration/backfill plan, reconciliation reports and rollback evidence.

Acceptance gate

- Target projects have complete registries and accountable historical records.
- Overrides and sources are preserved; unexplained losses or silently discarded files block acceptance.

Boundary and dependency

- Use verified backups and dry runs before applying controlled data changes.

Database-Driven Tab Migration

Move tabs to domain APIs without repeated extraction or private classification.

Work scope

- Migrate Documents, Drawings, Material Submittals, BOQ, Calculations, Compliance and Logs through explicit consumer contracts.
- Preserve each tab's domain behaviour and engineer controls; migrate Home after upstream consumers.
- Expose freshness, conflicts, missing data and processing state from stored records.

Deliverables

- Per-tab migration matrix, compatibility evidence, API contracts and integrated workflow tests.

Acceptance gate

- No extraction/model dispatch on ordinary read GETs.
- Each migrated tab uses the central truth and preserves existing functionality.

Boundary and dependency

- Migrate incrementally; remove a legacy producer only after its consumers have passed verification.

Stabilization, Dependency Engine & Legacy Cleanup

Make the central architecture stable, observable and safe under real workloads.

Work scope

- Finalize dependencies, retry rules, concurrency, stale-result prevention and interrupted-worker recovery.
- Measure performance and operational failure modes; remove duplicate legacy processing after migration.
- Plan SQLite-to-PostgreSQL migration once data ownership and workflows are stable, including transactions, locking, pooling and rollback.

Deliverables

- Operational runbooks, monitoring, recovery tests, cleanup record and PostgreSQL migration evidence when executed.

Acceptance gate

- No major duplicate processing remains; dependencies and workers recover safely.
- Production migration, if performed, reconciles all data and preserves functionality.

Boundary and dependency

- RAG, embeddings, fine-tuning and optional agent expansion follow M7 as separately scoped work; they are not extra numbered milestones.

Execution and evidence

Architecture

- File discovery -> document registry -> extraction/OCR -> classification -> central processors -> domain database -> APIs -> tabs.
- New or changed files are processed in background. Cached facts are reused; opening a tab should not repeat the complete PDF-reading cycle.

AI layer and PostgreSQL

- AI extraction/verification starts in M2; classification assistance belongs in M3. Evidence and domain contracts govern later usage.
- PostgreSQL belongs to stabilization after ownership/workflows settle. RAG and fine-tuning are later optional work, not prerequisites for extraction reliability.

Delegation recommendation

- Use the user-selected Fable orchestrator at Medium effort for coordination and bounded task assignment. Use Opus at High effort for difficult implementation and evidence interpretation.
- Use a fresh read-only reviewer for acceptance evidence. Verify available exact model IDs; do not silently substitute models or change frozen experiment pins.

Sources and status limits

- Primary: master-roadmap/MASTER-ROADMAP.md, Revision 01 (28 September 2026), and its AI Accuracy Policy.
- Status context: owner-supplied M2 history and declaration/verification reports. These planning PDFs do not authorize a live run or supersede frozen declarations.